

Bootstrapping

Simon Wood, University of Edinburgh, U.K.

Sampling distributions




- ▶ When we talk about the distribution of an estimator, $\hat{\theta}$, we are considering the randomness that $\hat{\theta}$ inherits from the randomness in the data \mathbf{y} .
- ▶ If we repeatedly replicated the process of gathering \mathbf{y} and estimating θ we would get a somewhat different $\hat{\theta}$ each time.
- ▶ This replicate to replicate variability is the *sampling distribution* of $\hat{\theta}$. The basis for inferences about θ .
- ▶ When considering maximum likelihood estimation, we looked at theoretical approximations for the distribution of $\hat{\theta}$.
- ▶ An alternative is to repeatedly simulate the data generation and estimation process to assess $\hat{\theta}$'s distribution. *Bootstrapping*.

Parametric bootstrapping

- ▶ Suppose we have a model specifying a p.d.f. $\pi_{\theta}(\mathbf{y})$ for our data.
- ▶ Let $\hat{\theta}$ be the MLE or other estimate of θ .
- ▶ Repeatedly simulate \mathbf{y}^* from $\pi_{\hat{\theta}}(\mathbf{y})$, and for each \mathbf{y}^* compute $\hat{\theta}^*$.
- ▶ The set of $\hat{\theta}^*$ are a *parametric bootstrap* approximation to the sampling distribution of $\hat{\theta}$.
- ▶ Typically the variability between the $\hat{\theta}^*$ is a reasonable approximation to the true sampling variability of $\hat{\theta}$.
- ▶ But the distribution of the bootstrap sample will be centred on $\hat{\theta}$, not the true θ .

Simple parametric bootstrap in R

```
mu <- mean(nhtemp) ## mean temperature
sd <- sd(nhtemp)   ## standard dev temp
n <- length(nhtemp) ## number of data
nb <- 10000        ## number of bootstrap samples
sdb <- mub <- rep(0,nb) ## replicate estimates
for (i in 1:nb) { ## parametric bootstrap loop
  yb <- rnorm(n,mu,sd) ## generate replicate data
  mub[i] <- mean(yb)   ## estimate mean
  sdb[i] <- sd(yb)     ## and standard deviation
}
```



Confidence intervals

- ▶ To illustrate that bootstrapping gets the variability right, compare the bootstrap standard deviation for $\hat{\mu}$ to its theoretical value.

```
> c(sd(mub), sd(nhtemp) / sqrt(n))  
[1] 0.1643081 0.1633892
```

- ▶ For more or less symmetrical sampling distributions *bootstrap percentile confidence intervals* can be computed.

```
> quantile(mub, c(.025, .975))  
 2.5%      97.5%  
50.83482 51.48179  
> quantile(sdb, c(.025, .975))  
 2.5%      97.5%  
1.042289 1.492834
```

- ▶ Obviously the parametric bootstrap also works with more complex models.

Generating random deviates in R

- ▶ R has built in functions for generating pseudorandom numbers from a wide variety of standard distributions.
- ▶ All these functions start with an 'r', take the number of deviates required as the first argument and further parameters of the distribution as further arguments.
- ▶ Examples are: `rnorm`, `rgamma`, `rpois`, `rbinom`, `rgeom`...
- ▶ All work by starting with uniform deviates* and transforming them. There are a variety of fast methods for doing this for specific distributions, but there is also a generic method:
 1. Generate a deviate, u , from $U(0, 1)$.
 2. Evaluate the inverse of the cumulative distribution function (CDF) for the distribution of interest (the *quantile function*) at u .
 3. The result is a random draw from the distribution of interest.

*see appendix C of *Core Statistics* for more on this.

CDFs and quantile functions

- ▶ For r.v. X , the CDF is defined as $F(x) = \Pr(X \leq x)$.
- ▶ The *quantile function* is the inverse of the CDF defined as $F^-(u) = \min(x|F(x) \geq u)$ (usual inverse for continuous F).
- ▶ If $X \sim F$ and $F(x)$ is continuous then $F(X) \sim U(0, 1)$.
Proof: Defining $U = F(X)$ and $u = F(x)$, we have




$$F(x) = \Pr(X \leq x) = \Pr(F(X) \leq F(x)) \Rightarrow u = \Pr(U \leq u)$$

but the r.h.s. is just the CDF of a $U(0, 1)$ r.v.

- ▶ Similarly $U \sim U(0, 1) \Rightarrow F^-(U) \sim F$, so we can generate from any distribution with a F^- function (continuous or not). In R
 - ▶ Quantile functions start with a 'q'. e.g. `qnorm`, `qpois` ...
So `qnorm(runif(3), 1, 2)` equivalent to `rnorm(3, 1, 2)`.
 - ▶ CDFs start with a 'p'. e.g. `pnorm`, `ppois`, `pgamma` ...

Example CDFs

- ▶ Here are three example CDFs for
 1. $N(0, 1)$.
 2. A mixture: 0.8 probability of drawing from $N(0, 1)$ and 0.2 probability of a zero.
 3. A Poisson with mean 3.



- ▶ Recall, $F(x) = \Pr(X \leq x)$ and $F^-(u) = \min(x | F(x) \geq u)$.


The empirical CDF

- ▶ Let x_1, x_2, \dots, x_n be sampled from a distribution with unknown CDF, $F(x)$.
- ▶ Since $F(x) = \Pr(X \leq x)$, we can estimate F by estimating the probabilities on the right, from our sample:

$$\hat{F}(x) = n^{-1} \sum_i \mathbb{I}(x_i \leq x)$$


where $\mathbb{I}(\cdot)$ is the indicator function (1/0 when argument is T/F).

- ▶ e.g. for $\mathbf{x} = (1, 0, 4.3, -.5, .8, 2.3, -.2, 1.1, 1.4, 2)$



Empirical CDF sampling: non-parametric bootstrapping

- ▶ Can generate from the estimated F , by generating $u \sim U(0, 1)$ and evaluating $\hat{F}^{-}(u)$ where $\hat{F}^{-}(u) = \min(x | \hat{F}(x) \geq u)$.




- ▶ ...but CDF steps are evenly spaced $1/n$ apart on the vertical axis and the jumps occur at the original x_i values...
- ▶ We are just randomly selecting from the original sample of x_i values with equal probability. *Non-parametric bootstrapping*.

Non-parametric bootstrap example

Suppose we are interested in the interquartile range of the New Haven temperature distribution.

```
nb <- 10000; n <- length(nhtemp)
iqb <- rep(0,nb) ## storage for bs interquartile range
for (i in 1:nb) { ## bootstrap loop
  yb <- sample(nhtemp,n,replace=TRUE) ## non-parametric bs sample
  iqb[i] <- diff(quantile(yb,c(.25,.75))) ## bootstrap iqr
}
## examine distribution (notice the skew)...
hist(iqb,freq=FALSE);lines(density(iqb,adjust=2))
```



Basic bootstrap confidence interval

- ▶ IQR distribution not at all symmetric. Percentile CI dubious.
- ▶ Really BS gives distribution of \hat{q}^* when true IQR is \hat{q} .
- ▶ Find interval containing 95% of \hat{q}^* . Write as $(\hat{q} - b, \hat{q} + c)$.

$$\begin{aligned} 0.95 &= \Pr(\hat{q} - b < \hat{q}^* < \hat{q} + c) = \Pr(-\hat{q}^* - b < -\hat{q} < -\hat{q}^* + c) \\ &= \Pr(\hat{q}^* + b > \hat{q} > \hat{q}^* - c) \end{aligned}$$

- ▶ View b and c as bootstrap estimates of constants such that

$$\Pr(\hat{q} + b > q > \hat{q} - c) = 0.95.$$

- ▶ That is $(\hat{q} - c, \hat{q} + b)$ is a 95% CI for q .
- ▶ Coincides with percentile interval for symmetric distribution.
- ▶ Obviously equally valid for parametric bootstrap.

Basic bootstrap CI R code

- ▶ Continuing the interquartile range example.

```
> iq <- diff(quantile(nhtemp,c(.25,.75))) ## original iqr
> ## 0.025 and 0.975 quantiles of b.s. irq...
> pci <- quantile(iqb,c(.025,.975))
> names(pci) <- names(iq) <- NULL ## avoid confusing names
> ## get upper and lower interval margins
> b <- iq - pci[1]; c <- pci[2] - iq
> c(iq-c,iq+b) ## basic CI
0.625 1.650
> pci ## equivalent percentile CI
1.000 2.025
```

- ▶ Notice the quite large difference between basic and percentile intervals in this case.

Bootstrapping multivariate data

- ▶ The bootstrapping idea generalizes readily to multivariate data.
- ▶ We simply resample *cases* - for tidy data this usually amounts to resampling rows. How?
- ▶ We can programme bootstrap resampling of univariate data by resampling either the data themselves, or their indices. That is

```
xb <- sample(x, n, replace=TRUE)
```

is equivalent to

```
bsi <- sample(1:n, n, replace=TRUE)
```

```
xb <- x[bsi]
```

- ▶ The second option generalizes immediately to resampling rows of a data matrix

```
bsi <- sample(1:n, n, replace=TRUE)
```


```
Xb <- X[bsi, ]
```

A more interesting example

- Consider a chemical experiment where you have measured yield, y , against flow rate of a reagent, x , and interest is in the x value maximizing yield.
- Suppose a quadratic model is appropriate

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \epsilon_i$$

where the β_j are parameters and ϵ_i is a zero mean r.v.



A more interesting bootstrap

- We can bootstrap to find a CI for the optimum yield. Suppose x and y data are in dataframe `dat`.

```
n <- nrow(dat); nb <- 1000; mx <- rep(0,nb)
for (i in 1:nb) { ## bootstrap loop
  ii <- sample(1:n,n,replace=TRUE) ## resample row indices
  b <- coef(lm(y ~ x + I(x^2),data=dat[ii,])) ## fit to resample
  mx[i] <- -b[2]/(2*b[3]) ## compute location of maximum
}
quantile(mx,c(0.05,0.95)) ## 90% CI
0.3785062 0.4458549
hist(mx)
```

