

Graphical models and automating Gibbs sampling

Simon Wood, University of Edinburgh, U.K.

Automatic Gibbs sampling

- ▶ Much of the pain can be removed from Gibbs sampling by automation.
- ▶ What makes that possible is the recognition that many Bayesian models can be abstractly represented as *Directed Acyclic Graphs*.
- ▶ Mathematically a graph* is a set of nodes connected in some way by edges. In our case
 - ▶ Nodes are variables or constants, such as data, model parameters and fixed parameters of priors.
 - ▶ Edges show dependencies between nodes.
- ▶ In a *directed* graph edges have direction, shown by arrows, with the node at the pointy end being the *child* of the *parent* at the other end. Parents directly control children.



* not to be confused with a *plot*

Types of node and edge

- ▶ Nodes and edges are of different types.
- ▶ A node with no parents is a *constant*
- ▶ A continuous arrow denotes a *stochastic dependence* – the distribution of the child node depends on the parent node. In this case the child node is a *stochastic* node. e.g.



... constant parent, stochastic child.


- ▶ A dashed arrow denotes a *deterministic dependence* – the child depends deterministically (not randomly) on the parent, and is a *deterministic* node. e.g.



... the child is a deterministic node.

The simple example again

- ▶ Recall the model $x_i \sim N(\mu, \sigma^2)$, with priors
 - ▶ $\tau = 1/\sigma^2 \sim \text{gamma}(a, b)$, i.e. prior $\pi(\tau) = b^a \tau^{a-1} e^{-b\tau} / \Gamma(a)$
 - ▶ Independently, $\mu \sim N(c, d)$.
- ▶ Its graphical representation, showing just 3 of the x_i nodes, is



- ▶ The graph is a directed *acyclic* graph (DAG), because, following the arrows, no path ever revisits a node.
- ▶ That the x_i are independent draws from $N(\mu, \sigma^2)$ is encoded in the lack of direct edges between them.

Why the DAG is useful...

- ▶ Let z_i denote the variable corresponding to the i^{th} node of the graph. Given the DAG structure of the model, we can write the joint density over all non-constant nodes as

$$\pi(\mathbf{z}) = \prod_i \pi(z_i | \text{parent}\{z_i\})$$

- ▶ Then, from the definition of a conditional p.d.f.

$$\pi(z_j | \mathbf{z}_{-j}) = \frac{\pi(\mathbf{z})}{\int \pi(\mathbf{z}) d\mathbf{z}_j} = \frac{\prod_i \pi(z_i | \text{parent}\{z_i\})}{\int \prod_i \pi(z_i | \text{parent}\{z_i\}) d\mathbf{z}_j},$$

- ▶ But only z_j dependent terms stay in the integral, the rest cancel

$$\begin{aligned} \pi(z_j | \mathbf{z}_{-j}) &= \frac{\pi(z_j | \text{parent}\{z_j\}) \prod_{i \in \text{child}\{j\}} \pi(z_i | \text{parent}\{z_i\})}{\int \pi(z_j | \text{parent}\{z_j\}) \prod_{i \in \text{child}\{j\}} \pi(z_i | \text{parent}\{z_i\}) d\mathbf{z}_j} \\ &\propto \pi(z_j | \text{parent}\{z_j\}) \prod_{i \in \text{child}\{j\}} \pi(z_i | \text{parent}\{z_i\}), \end{aligned}$$

... why the DAG is useful

- ▶ So however complicated the model, and however large its DAG, the conditional density of node z_j depends only on its immediate ‘family’ — its parents, children, and ‘partners’ (other parents of its children).
- ▶ The Gibbs update of z_j only needs to know the state of those ‘family’ nodes.
- ▶ The fact that the DAG allows the conditionals to be modularised in this way:
 1. makes computation efficient as the conditionals only require evaluations over a small portion of the graph.
 2. facilitates automation of the process of identifying conditionals (using known results on *conjugacy* of distributions).
- ▶ JAGS (Just Another Gibbs Sampler) is a stand alone software package for automatic Gibbs sampling. `rjags` interfaces to R.

JAGS basic use

- ▶ JAGS requires that your model is specified in the JAGS language written down in a text file.
- ▶ This text file should be in the working directory, which can be checked by `getwd` and set by `setwd`.
- ▶ From within R the file is then compiled into a Gibbs sampler via a call to the function `jags.model`.
- ▶ The returned sampler object can then be used to simulate from the model posterior using the `jags.samples` function.
- ▶ An alternative to `jags.samples` is `coda.samples` which returns an object of class `mcmc.list` suitable for direct use with the `coda` package for checking and post-processing MCMC output.
- ▶ Model comparison via the *deviance information criterion* (DIC) is facilitated by `dic.samples` (requires the model to have been compiled by `jags.models` with at least `n.chains=2`).

Useful coda functions

- ▶ `plot.mcmc.list` plot method for MCMC simulation output.
- ▶ `acfplot` plots chain ACFs compactly.
- ▶ `autocorr` and `crosscorr` for examining within chain correlation.
- ▶ `effectiveSize` to get the effective sample sizes of simulation output.
- ▶ `HPDinterval` for highest posterior density intervals.