

Gibbs sampling for Bayesian Inference

Simon Wood, University of Edinburgh, U.K.

One at a time updating

- ▶ In MH sampling it can be difficult to find a distribution that makes good proposals for the whole parameter vector, θ .
- ▶ But there is nothing in the MH scheme to prevent us using a proposal that updates the parameters in smaller blocks, or even one at a time, cycling through the elements of θ .
- ▶ Such single element proposals are easier to tune.
- ▶ Another interesting fact is that if our proposal distribution is $\pi(\theta|\mathbf{y})$ then the MH acceptance probability is exactly 1 (try it).
- ▶ This observation is not very practical for all at once updating of θ , but it turns out to be very useful for one at a time updating, and is one way of motivating *Gibbs sampling*.

Gibbs sampling

- ▶ Let $\boldsymbol{\theta}_{-i}$ denote the elements of $\boldsymbol{\theta}$ other than θ_i .
- ▶ Suppose we want a random draw from $\pi(\boldsymbol{\theta}|\mathbf{y})$, and we already have a random draw, $\boldsymbol{\theta}_{-i}^*$, from $\pi(\boldsymbol{\theta}_{-i}|\mathbf{y})$.
- ▶ All we need to do is generate θ_i^* from $\pi(\theta_i|\boldsymbol{\theta}_{-i}^*, \mathbf{y})$ and append this to $\boldsymbol{\theta}_{-i}^*$ to get $\boldsymbol{\theta}^* \sim \pi(\boldsymbol{\theta}|\mathbf{y}) = \pi(\theta_i|\boldsymbol{\theta}_{-i}^*, \mathbf{y})\pi(\boldsymbol{\theta}_{-i}|\mathbf{y})$.
- ▶ To get a slightly different draw from the same distribution we could discard θ_i^* and draw a fresh value in the same way.
- ▶ We could repeatedly cycle through all elements of $\boldsymbol{\theta}$, dropping each and replacing with a draw from its conditional $\pi(\theta_i|\boldsymbol{\theta}_{-i}^*, \mathbf{y})$.
- ▶ This method is *Gibbs sampling*.
- ▶ To use it we need to identify the conditional densities.
- ▶ When identification of a conditional is not possible, we can use a Metropolis Hastings step for that parameter, or more efficient one dimensional alternatives.

A simple example

- ▶ Consider n observations from model $x_i \sim N(\mu, \sigma^2)$, with priors
 - ▶ $\tau = 1/\sigma^2 \sim \text{gamma}(a, b)$, i.e. prior $\pi(\tau) = b^a \tau^{a-1} e^{-b\tau} / \Gamma(a)$
 - ▶ Independently, $\mu \sim N(c, d)$.
- ▶ The joint density of \mathbf{x} , μ and τ is

$$\begin{aligned}\pi(\mathbf{x}, \tau, \mu) &= \pi(\mathbf{x}|\tau, \mu)\pi(\tau)\pi(\mu) \\ &\propto \tau^{n/2} e^{-\sum_i \tau(x_i - \mu)^2/2} e^{-(\mu - c)^2/(2d)} \tau^{a-1} e^{-b\tau}\end{aligned}$$


- ▶ By Bayes Theorem, the conditional for each parameter will be proportional to $\pi(\mathbf{x}, \tau, \mu)$. The trick is to recognize the conditional from its kernel - i.e from the terms actually involving the parameter.
- ▶ e.g. $\pi(\tau|\mathbf{x}, \mu) \propto \tau^{n/2+a-1} e^{-\sum_i \tau(x_i - \mu)^2/2 - b\tau}$, recognizable as a $\text{gamma}(n/2 + a, \sum_i (x_i - \mu)^2/2 + b)$.
- ▶ Similarly, but with a bit more effort,

$$\mu|\mathbf{x}, \tau \sim N\{(dn\bar{x}\tau + c)/(dn\tau + 1), d/(dn\tau + 1)\}$$

Checking the priors

- ▶ Let's apply the preceding model to the `nhtemp` data. This will involve choosing the constants a , b , c and d defining the priors.
- ▶ We should look at the data, and check the priors...

```
x <- nhtemp ## just to save later typing
a <- 1.2; b <- .6 ## gamma prior shape and scale
c <- 50; d <- 100 ## normal prior mean and variance
## check data and priors...
par(mfrow=c(1,3),mar=c(4,4,1,1))
hist(x,main="")
tau <- seq(0,10,length=400)
plot(tau,dgamma(tau,a,b),type="l")
mu <- seq(0,100,length=400)
plot(mu,dnorm(mu,c,d^.5),type="l")
```



Gibbs sampling loop

- Now we can run the Gibbs sampling loop...

```
ns <- 10000 ## number of samples
th <- matrix(0,2,ns) ## sample storage
mu <- 0;tau <- 0 ## initial states
n <- length(x)
## store constants needed repeatedly...
dn <- d*n; dnx <- dn*mean(x)


for (i in 1:ns) { ## Gibbs loop
  ## update mu | tau, x ...
  mu <- rnorm(1, (dnx*tau+c)/(dn*tau+1), sqrt(d/(dn*tau+1)))
  ## update tau | mu, x ...
  tau <- rgamma(1, n/2+a, sum((x-mu)^2)/2 + b)

  th[,i] <- c(mu, 1/sqrt(tau)) ## store as mean and sd
}
```

Checking the chains

- First look at the trace plots...

```
par(mfrow=c(2,1))  
## see ?plotmath for adding maths to plots  
plot(th[1,],type="l",ylab=expression(mu))  
plot(th[2,],type="l",ylab=expression(sigma))
```




- Convergence appears to be very rapid, and mixing good.

Checking autocorrelation

- Let's check the impression of rapid mixing with ACF plots...

```
par(mfrow=c(1,2))  
acf(th[1,]); acf(th[2,])
```





- The samples appear to be effectively independent of each other - actually this behaviour is special to this example - parameters with high posterior correlation would give chains with higher autocorrelation.

Did the data matter?

- ▶ A common error is to believe that having updated your beliefs about the parameters using data, then *all* those beliefs are now validated by data. The Covid literature provides many examples.
- ▶ To avoid this it is a good idea to compare priors to posteriors, to see what effect the data had, if any.

```
hist(th[1,],xlab=expression(mu),main="",probability=TRUE)
mu <- seq(50,52,length=100)
lines(mu,dnorm(mu,c,d^.5),col=2)
hist(1/th[2,]^2,xlab=expression(tau),main="",probability=TRUE)
tau <- seq(0.1,1.2,length=100)
lines(tau,dgamma(tau,a,b),col=2)
```



- ▶ The posteriors are very different to the priors, which have little influence - here the posteriors are determined by the data.

Posterior means and CIs

- ▶ Nothing is different here to in the MH sampling case (except that we only have 2 parameter in this example).

```
> pm <- rowMeans(th)
> names(pm) <- c("mu", "sigma")
> pm ## posterior mean
      mu      sigma
51.158764  1.263793
>
> ci <- apply(th, 1, quantile, probs=c(.025, .975))
> colnames(ci) <- c("mu", "sigma")
> ci ## confidence intervals
      mu      sigma
2.5%  50.83651  1.064045
97.5%  51.48939  1.511449
```